

Introduction

In this assignment I wrote a Map-Reduce program to compute analytics (max., min., median, standard deviation) on a public data set regarding Airborne Radioactivity. The program is written in Scala and executed in the Apache Spark cluster-computing framework.

Requirements

- a) Compute Maximum, Minimum, Median and Standard Deviation of the metric “7Be MDC/7Be CMD (mBq/m3)” from the “The Canadian Radiological Monitoring Network – Airborne Radioactivity” data set (csv).
- b) Display analytics of the above-mentioned metric for a location in each year.
- c) The software is programmed and executed in a Map-Reduce runtime such as Apache Hadoop or AWS EMR. It cannot be a standalone program mimicking map-reduce behavior.
- d) The execution environment is either a local computer or on the Cloud. Submission in a single package with source code and executable.

Specifications

- a) Median is estimated as the 50th percentile using “percentile_approx”. Standard deviation is based on the “stddev_samp” (square root of variance of the sample). Min and Max is straightforward.
- b) The data set defines location as city, and the date is in MM/DD/YYYY format so we need to extract the year as a substring from the date column. The city and year form the fields on which we group to calculate the aggregates.
- c) Programming language used is SCALA and the Map-Reduce framework is Apache Spark.
- d) Execution environment is local computer using a standalone Spark cluster, with a master and workers running locally. The results are combined and outputted in “results.csv” for every location/year pair.

Implementation

a) Map-Reduce Algorithm Design

In the Apache Spark framework there are no explicit “map” and “reduce” functions to define. Hence, I jumped into the implementation:

First, the CSV file is loaded (lazy) within an SQL context using the spark-csv package:

```
val full_csv = sqlContext.read
  .format("com.databricks.spark.csv")
  .option("header", "true") // Use first line of all files as header
  .option("inferSchema", "true") // Automatically infer data types (otherwise
  everything is assumed string)
  .load(csvFile)
```

The next step is to select the required fields for our analytics based on the requirements:

```
val data = full_csv.select("Location/Emplacement",
                           "Collection Start/Debut du prelevement (UTC)",
                           "7Be MDC/7Be CMD (mBq/m3)")
```

From the data set, it was clear the “date” is not properly formatted (MM/DD/YYYY) and so the year (YYYY) is extracted from the date and the columns are renamed for the sake of clarity:

```
val newNames = Seq("location", "date", "mdc")
val df = data.toDF(newNames: _*).withColumn("year", substring_index(col("date"),
"/", -1))
```

The required Analytics can now be performed on this “cleaned” data set. The operations we want are “min”, “max”, “stddev” and “median” of the MDC metric for all locations per year:

```
val mapping: Map[String, Column => Column] = Map(
  "min" -> min, "max" -> max, "mean" -> avg, "stddev" -> stddev)

val groupBy = Seq("location", "year")
val aggregate = Seq("mdc")
val operations = Seq("min", "max", "mean", "stddev")
val exprs = aggregate.flatMap(c => operations .map(f => mapping(f)(col(c))))

df.registerTempTable("df ")
var median = sqlContext.sql("select location, year, percentile_approx(mdc,
0.5) as median from df group by location, year")
```

This is the core of our “map-reduce” algorithm design, we group the data by “location” and then by “year” and perform the aggregate operations “min”, “max”, “mean” (extra), “stddev” on “mdc”. These operations are already defined and optimized within the sql.functions package.

Since the median function does not exist, we write the sql query based on the ‘percentile_approx’ function to approximate the 50th percentile.

```
val results = df.groupBy(groupBy.map(col): _*).agg(exprs.head, exprs.tail: _*)  
  
val end_result = results.join(median, Seq("location", "year"), joinType="outer")
```

Essentially, the “MAP” and “REDUCE” portions of our algorithm are handled internally by the Spark framework. I do not explicitly define them, rather, I make use of pre-defined grouping and functions on the data frame.

To get the results, we execute the above defined operations on our dataset and join the previous results with the median results.

```
end_result.repartition(1).select("location", "year", "min(mdc)", "max(mdc)", "avg(mdc)", "stddev_samp(mdc)", "median")  
    .write.format("com.databricks.spark.csv")  
    .option("header", "true")  
    .save("results.csv")
```

Finally we repartition (merge) the results together into one csv file and save it in the “results.csv” folder.

b) HOW-TO Run the Program

Requirements:

Java v1.7+
Scala v2.1+
Apache Spark binaries
Hadoop's winutils.exe tool
Environment variables
Windows 7

Installation guide and links for the required binaries can be found here:

<https://edumine.wordpress.com/2015/06/11/how-to-install-apache-spark-on-a-windows-7-environment/>

Running the program:


1) **Run the Master** with following command:

```
spark-class org.apache.spark.deploy.master.Master
```

If successful, you should see:

```
17/11/25 15:31:25 INFO Utils: Successfully started service 'sparkMaster' on port 7077.
17/11/25 15:31:25 INFO Master: Starting Spark master at spark://192.168.2.13:7077
17/11/25 15:31:25 INFO Master: Running Spark version 2.2.0
17/11/25 15:31:25 INFO Utils: Successfully started service 'MasterUI' on port 8080.
17/11/25 15:31:26 INFO MasterWebUI: Bound MasterWebUI to 0.0.0.0, and started at http://192.168.2.13:8080
17/11/25 15:31:26 INFO Utils: Successfully started service on port 6066.
17/11/25 15:31:26 INFO StandaloneRestServer: Started REST server for submitting applications on port 6066
17/11/25 15:31:26 INFO Master: I have been elected leader! New state: ALIVE
```

The master also has a default webpage visible at “localhost:8080”:



URL: spark://192.168.2.13:7077
REST URL: spark://192.168.2.13:6066 (*cluster mode*)
Alive Workers: 1
Cores in use: 6 Total, 0 Used
Memory in use: 7.0 GB Total, 0.0 B Used
Applications: 0 Running, 3 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers

| Worker Id | Address | State | Cores | Memory |
|---|-------------------|-------|------------|---------------------|
| worker-20171125153655-192.168.2.13-6018 | 192.168.2.13:6018 | ALIVE | 6 (0 Used) | 7.0 GB (0.0 B Used) |

Running Applications

| Application ID | Name | Cores | Memory per Executor | Submitted Time | User | State | Duration |
|----------------|------|-------|---------------------|----------------|------|-------|----------|
|----------------|------|-------|---------------------|----------------|------|-------|----------|

2) **Run at-least ONE Worker** under the master with the following command:

```
spark-class org.apache.spark.deploy.worker.Worker spark://192.168.2.13:7077
```

Notice the worker is targeted to where the Spark master has been ran, i.e. “spark://192.168.2.13:7077”.

In this step, we can also specify the amount of memory and cores the worker is limited to.

If successful, the worker will be registered with the master:

```
17/11/25 15:36:55 INFO Utils: Successfully started service 'sparkWorker' on port 6018.
17/11/25 15:36:56 INFO Worker: Starting Spark worker 192.168.2.13:6018 with 6 cores, 7.0 GB RAM
17/11/25 15:36:56 INFO Worker: Running Spark version 2.2.0
17/11/25 15:36:56 INFO Worker: Spark home: C:\spark-2.2.0-bin-hadoop2.7\bin\..
17/11/25 15:36:56 INFO Utils: Successfully started service 'WorkerUI' on port 8081.
17/11/25 15:36:56 INFO WorkerWebUI: Bound WorkerWebUI to 0.0.0.0, and started at http://192.168.2.13:8081
17/11/25 15:36:56 INFO Worker: Connecting to master 192.168.2.13:7077...
17/11/25 15:36:56 INFO TransportClientFactory: Successfully created connection to 192.168.2.13:7077 after 70 ms (0 ms spent in bootstraps)
17/11/25 15:36:56 INFO Worker: Successfully registered with master spark://192.168.2.13:7077
```

And from the master’s perspective:

```
17/11/25 15:36:56 INFO Master: Registering worker 192.168.2.13:6018 with 6 cores, 7.0 GB RAM
```

3) **Submit the job** to the cluster with the following command (within project directory):

```
spark-submit target/scala-2.10/simple-project_2.10-1.0.jar spark://192.168.2.13:7077 \ --class MapReduce
```

Optional: Recompile the package by using: **sbt package**

On the master's UI, you can see the job:

While it's running:


Running Applications


| Application ID | Name | Cores | Memory per Executor | Submitted Time | User | State | Duration |
|---|-------------------------|-------|---------------------|---------------------|------|---------|----------|
| app-20171125163554-0003 (kill) | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:35:54 | Jaw | RUNNING | 4 s |

And after completion:

Completed Applications

| Application ID | Name | Cores | Memory per Executor | Submitted Time | User | State | Duration |
|---|-------------------------|-------|---------------------|---------------------|------|----------|----------|
| app-20171125161636-0002 | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:16:36 | Jaw | FINISHED | 25 s |

4) **View the results** inside the  **results.csv** folder as:

 [part-00000-8775cc8d-a48c-48b5-bf00-24...](#)

This is the CSV file with all the parts merged. You can now sort it by 'location' then 'year' to view the results neatly:

| location | year | min(mdc) | max(mdc) | avg(mdc) | stddev_sa | median |
|----------|------|----------|----------|----------|-----------|----------|
| Alert | 2009 | 0.123499 | 0.482821 | 0.236398 | 0.086713 | 0.210171 |
| Alert | 2010 | 0.112835 | 3.023419 | 1.010645 | 0.703639 | 1.008268 |
| Alert | 2011 | 0.193144 | 0.965462 | 0.411703 | 0.2093 | 0.320574 |
| Alert | 2012 | 0.187856 | 0.787936 | 0.430048 | 0.151338 | 0.404686 |
| Alert | 2013 | 0.163912 | 0.976627 | 0.414148 | 0.199903 | 0.355648 |
| Alert | 2014 | 0.13779 | 0.739147 | 0.338335 | 0.131382 | 0.316775 |
| Alert | 2015 | 0.171283 | 0.564321 | 0.333889 | 0.109639 | 0.308937 |
| Alert | 2016 | 0.211878 | 0.566544 | 0.349633 | 0.086807 | 0.329935 |
| Alert | 2017 | 0.179611 | 0.203982 | 0.192603 | 0.012265 | 0.194218 |

c) Timing of jobs

Unfortunately Apache Spark does not provide timings in the logs for the map task and reduce task separately.

However, we have access to the combined time for the job which seems to actually vary. Here's some sample timings:

Completed Applications

| Application ID | Name | Cores | Memory per Executor | Submitted Time | User | State | Duration |
|---|-------------------------|-------|---------------------|---------------------|------|----------|----------|
| app-20171125163554-0003 | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:35:54 | Jaw | FINISHED | 14 s |
| app-20171125161636-0002 | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:16:36 | Jaw | FINISHED | 25 s |
| app-20171125161124-0001 | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:11:24 | Jaw | FINISHED | 18 s |
| app-20171125161058-0000 | COEN 424 - Assignment 2 | 6 | 1024.0 MB | 2017/11/25 16:10:58 | Jaw | FINISHED | 20 s |

d) Data partition and workload balancing

I tested having 3 workers in total doing the same job compared to 1 worker:

Workers

| Worker Id | Address | State | Cores | Memory |
|---|-------------------|-------|------------|-------------------------|
| worker-20171125153655-192.168.2.13-6018 | 192.168.2.13:6018 | ALIVE | 6 (6 Used) | 7.0 GB (1024.0 MB Used) |
| worker-20171125164646-192.168.2.13-6736 | 192.168.2.13:6736 | ALIVE | 6 (6 Used) | 7.0 GB (1024.0 MB Used) |
| worker-20171125164659-192.168.2.13-6768 | 192.168.2.13:6768 | ALIVE | 6 (6 Used) | 7.0 GB (1024.0 MB Used) |

Running Applications

| Application ID | Name | Cores | Memory per Executor | Submitted Time | User | State | Duration |
|-----------------------------------|-------------------------|-------|---------------------|---------------------|------|---------|----------|
| app-20171125164714-0004 (kill) | COEN 424 - Assignment 2 | 18 | 1024.0 MB | 2017/11/25 16:47:14 | Jaw | RUNNING | 3 s |

Each worker gets assigned up to 228 tasks:

```
17/11/25 16:36:08 INFO CoarseGrainedExecutorBackend: Got assigned task 228
```

With 3 workers it's about the same number of tasks per worker:

Worker 1: `17/11/25 16:47:45 INFO CoarseGrainedExecutorBackend: Got assigned task 224`

Worker 2: `17/11/25 16:47:45 INFO CoarseGrainedExecutorBackend: Got assigned task 227`

Worker 3: `17/11/25 16:47:45 INFO CoarseGrainedExecutorBackend: Got assigned task 228`

However, **the difference is the size of each task:**

For 1 worker:

```
17/11/25 16:36:08 INFO MemoryStore: Block broadcast_11_piece0 stored as bytes in memory (estimated size 25.7 KB, free 364.6 MB)
Finished task 0.0 in stage 9.0 (TID 205). 4013 bytes result sent to driver
```

For 3 workers:

```
17/11/25 16:47:36 INFO MemoryStore: Block broadcast_7_piece0 stored as bytes in memory (estimated size 11.2 KB, free 366.3 MB)
17/11/25 16:47:38 INFO Executor: Finished task 0.0 in stage 3.0 (TID 3). 2883 bytes result sent to driver
```

Another thing I noticed is with 3 workers, since the partitions are smaller

Therefore, load balancing is done via data partitioning. Although the numbers of chunks are the same for a different number of workers, each worker actually gets a smaller data chunks the more workers exist.

Furthermore, I noticed the 3 workers more quickly and frequently finishing their tasks since their tasks are actually smaller.

For 3 workers:

```
17/11/25 16:47:39 INFO Executor: Finished task 6.0 i
17/11/25 16:47:39 INFO Executor: Finished task 23.0
17/11/25 16:47:39 INFO Executor: Finished task 26.0
17/11/25 16:47:39 INFO Executor: Finished task 16.0
```

Notice the tasks are being quickly finished. For the case with 1 worker, there's more time between tasks.

This load balancing is also dependent on the resources of each worker. In my case all workers had the same amount of memory and processing power. However, the master would allocate properly sized chunks for a slower worker.

Sample results:

| | | | | | | |
|------------|------|----------|----------|----------|----------|----------|
| Winnipeg | 2017 | 0.012894 | 0.039385 | 0.028413 | 0.007981 | 0.029877 |
| Inuvik | 2013 | 0.02306 | 0.04583 | 0.034654 | 0.006949 | 0.03287 |
| Saskatoon | 2014 | 0.021946 | 0.027972 | 0.025215 | 0.001929 | 0.025051 |
| Winnipeg | 2013 | 0.01779 | 0.059325 | 0.035264 | 0.008992 | 0.035654 |
| Quebec Ci | 2015 | 0.013091 | 0.054384 | 0.03002 | 0.009324 | 0.026224 |
| Toronto | 2015 | 0.019737 | 0.05841 | 0.03253 | 0.008529 | 0.02909 |
| Halifax | 2014 | 0.010372 | 0.045545 | 0.031647 | 0.008075 | 0.035442 |
| Inuvik | 2016 | 0.022501 | 0.047701 | 0.033088 | 0.008061 | 0.028568 |
| Digby | 2016 | 0.016888 | 0.046437 | 0.029059 | 0.007373 | 0.027271 |
| Moncton | 2015 | 0.022086 | 0.505067 | 0.041737 | 0.073601 | 0.026687 |
| Moosonee | 2015 | 0.023146 | 0.046816 | 0.031924 | 0.0075 | 0.027804 |
| Quebec Ci | 2010 | 0.021949 | 0.053726 | 0.032094 | 0.006797 | 0.029634 |
| Iqaluit | 2016 | 0.004202 | 0.03473 | 0.006345 | 0.00423 | 0.00521 |
| Coral Harb | 2015 | 0.005483 | 0.146528 | 0.041899 | 0.032665 | 0.03046 |
| RPB | 2017 | 0.019757 | 0.041683 | 0.030095 | 0.007895 | 0.026212 |
| Toronto | 2009 | 0.020999 | 0.070343 | 0.031878 | 0.01104 | 0.028844 |
| Coral Harb | 2016 | 0.010692 | 0.137449 | 0.038933 | 0.026334 | 0.032397 |
| Moncton | 2016 | 0.02256 | 0.047133 | 0.031985 | 0.007679 | 0.02788 |
| Moncton | 2017 | 0.024259 | 0.05887 | 0.038728 | 0.012527 | 0.042818 |
| Montreal | 2009 | 0.019071 | 0.059454 | 0.031411 | 0.008417 | 0.030154 |
| Montreal | 2010 | 0.02038 | 0.045422 | 0.031716 | 0.006705 | 0.029695 |
| Montreal | 2011 | 0.019943 | 0.047609 | 0.030869 | 0.007409 | 0.029634 |
| Montreal | 2012 | 0.019037 | 0.050955 | 0.031824 | 0.0078 | 0.028921 |
| Montreal | 2013 | 0.020138 | 0.068425 | 0.032352 | 0.009191 | 0.030538 |
| Montreal | 2014 | 0.019677 | 0.047425 | 0.030626 | 0.007324 | 0.028048 |
| Montreal | 2015 | 0.019873 | 0.046298 | 0.031198 | 0.007465 | 0.027573 |
| Montreal | 2016 | 0.018379 | 0.056021 | 0.032485 | 0.008759 | 0.031539 |
| Montreal | 2017 | 0.012971 | 0.054468 | 0.034181 | 0.012614 | 0.031975 |
| Moosonee | 2009 | 0.019639 | 0.064352 | 0.030452 | 0.009134 | 0.026125 |
| Moosonee | 2010 | 0.022226 | 0.046052 | 0.033568 | 0.006702 | 0.032918 |
| Moosonee | 2011 | 0.021824 | 0.053053 | 0.03313 | 0.007738 | 0.03133 |
| Moosonee | 2012 | 0.02152 | 0.054256 | 0.032921 | 0.008213 | 0.030138 |
| Moosonee | 2013 | 0.0199 | 0.147234 | 0.034571 | 0.017765 | 0.029288 |
| Moosonee | 2014 | 0.022184 | 0.045323 | 0.032216 | 0.007174 | 0.029782 |
| Moosonee | 2015 | 0.023146 | 0.046816 | 0.031924 | 0.0075 | 0.027804 |
| Moosonee | 2016 | 0.023263 | 0.070268 | 0.033363 | 0.009454 | 0.029269 |
| Moosonee | 2017 | 0.0249 | 0.047734 | 0.034524 | 0.009585 | 0.028538 |
| Ottawa | 2009 | 0.019449 | 0.065883 | 0.030208 | 0.008781 | 0.026945 |
| Ottawa | 2010 | 0.014308 | 0.044299 | 0.02984 | 0.006897 | 0.028568 |

Full results inside the .csv file within the "results.csv" directory.